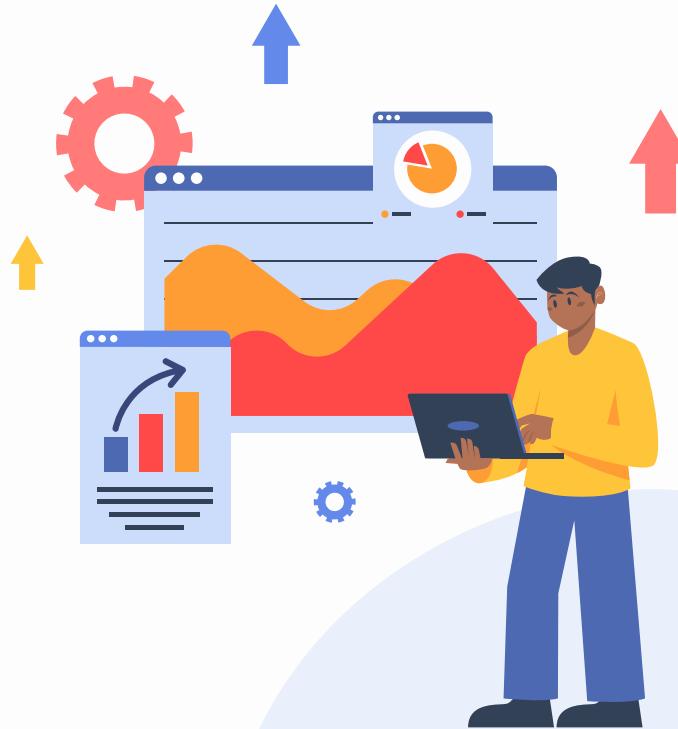


# Doordash Prediction Project

By: Abigail Chen, Simran Kaur, Xincheng Li,  
Sofia Luong





# Problem Scope & Data

# Problem:

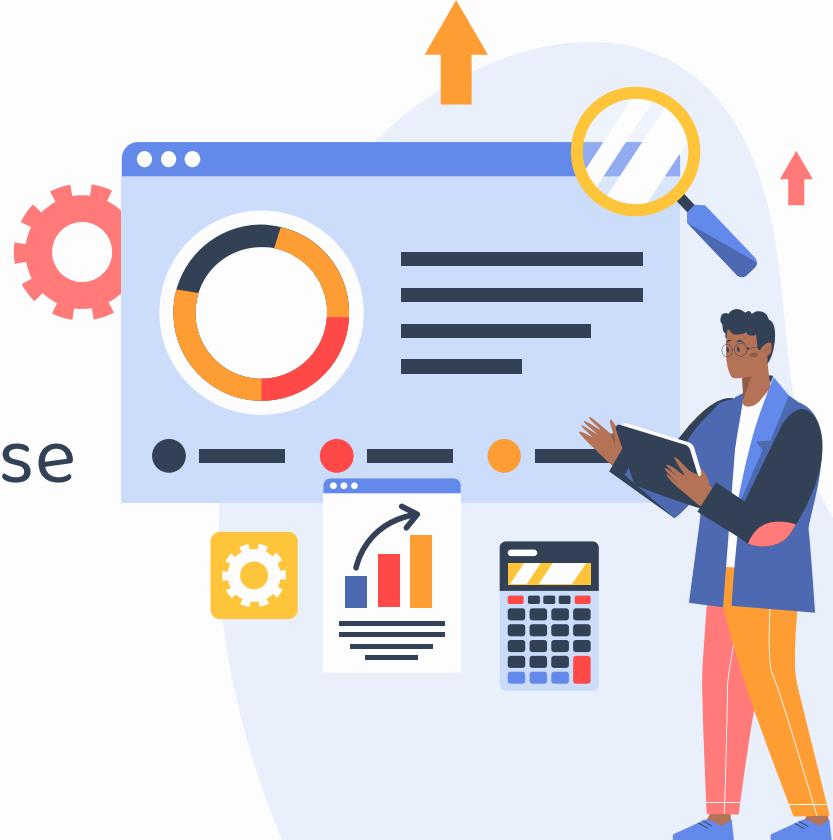
Optimizing the prediction of delivery time

DoorDash ETA Prediction



# Motivation:

Help users gain more precise information on their deliveries





# Table of contents

**01 Data cleaning,  
Pre-Processing and EDA**

**02 Unsupervised  
Learning**

**03 Supervised Learning  
and Key Methodology**

**04 Presentation of  
Results**





# 01

# Data cleaning, Pre-Processing and EDA

# Steps We Took for Preprocessing/Data Cleaning



**Feature  
Creation**



**Dealing with  
NaN Values**



**Removing  
Features**



**Removing  
Outliers**



**Normalizing  
Data**





# Feature Engineering - Feature Creation

- **Purpose:**
  - 'created\_at' & 'actual\_delivery\_time' columns
  - No column for total delivery time (our response variable)
- **Approach:**
  - 'Total\_seconds' = 'actual\_delivery\_time' - 'created\_at'
  - 'Weekday' and 'hour\_of\_day' feature using 'created\_at'
  - Dropped 'created\_at' and 'actual\_delivery\_time' features



# Columns used for Feature Creation

	market_id	created_at	actual_delivery_time	store_id	store_primary_category	order_protocol	total_items	subtotal	num_distinct_items	min_item_price
0	1.0	2015-02-06 22:24:17	2015-02-06 23:27:16	1845	american	1.0	4	3441	4	557
1	2.0	2015-02-10 21:49:25	2015-02-10 22:56:29	5477	mexican	2.0	1	1900	1	1400
2	3.0	2015-01-22 20:39:28	2015-01-22 21:09:09	5477	Nan	1.0	1	1900	1	1900
3	3.0	2015-02-03 21:21:45	2015-02-03 22:13:00	5477	Nan	1.0	6	6900	5	600
4	3.0	2015-02-15 02:40:36	2015-02-15 03:20:26	5477	Nan	1.0	3	3900	3	1100



# Code

```
data_cleaned = data.copy()
created_at = pd.to_datetime(data_cleaned['created_at'])
actual_delivery_time = pd.to_datetime(data_cleaned['actual_delivery_time'])

data_cleaned['hour_of_day'] = created_at.dt.hour
data_cleaned['weekday'] = actual_delivery_time.dt.weekday

data_cleaned['total_seconds'] = (actual_delivery_time - created_at).dt.total_seconds()
data_cleaned = data_cleaned[data_cleaned['total_seconds']>0]

data_cleaned = data_cleaned.drop(["actual_delivery_time", "created_at"], axis=1)
```



# Dealing with Null values

- **Approach 1 → Dropped rows with missing values:**
  - Columns with large number of missing values
    - 'total\_outstanding\_orders', 'total\_busy\_dashers', 'total\_onshift\_dashers'
  - Columns with discrete values
    - 'market\_id' and 'order\_protocol'
- **Approach 2 → Imputed missing values with mean:**
  - Columns with few missing values were imputed with their means.



# Null Values

market_id	987
store_id	0
store_primary_category	4760
order_protocol	995
total_items	0
subtotal	0
num_distinct_items	0
min_item_price	0
max_item_price	0
total_onshift_dashers	16262
total_busy_dashers	16262
total_outstanding_orders	16262
estimated_order_place_duration	0
estimated_store_to_consumer_driving_duration	526
hour_of_day	0
weekday	0
total_seconds	0

Too many missing values



# Code

```
data_cleaned2 = data_cleaned.dropna(subset=['total_outstanding_orders', 'total_busy_dashers', 'order_protocol',
                                             'total_onshift_dashers', 'market_id'])

data_cleaned3 = data_cleaned2.drop(['store_id', 'store_primary_category'], axis=1)

data_cleaned3 = data_cleaned3.fillna(data_cleaned3.mean())
```



# Feature Engineering - Removing Features

- 'store\_primary\_category' was categorical.
  - Too many unique values to one-hot encode
  - low correlation with 'total\_seconds'
- Removed features with extremely high correlations with other predictors
  - 'num\_distinct\_items'
  - 'total\_onshift\_dashers'
  - 'total\_busy\_dashers'
  - 'order\_protocol'



# Correlation Matrix

	market_id	order_protocol	total_items	subtotal	num_distinct_items	min_item_price	max_item_price	total_onshift_dashers
market_id	1.000000	-0.014876	0.002115	-0.002125	0.013303	-0.009192	-0.005576	0.069461
order_protocol	-0.014876	1.000000	0.005411	-0.052642	-0.026407	-0.041685	-0.088698	0.145208
total_items	0.002115	0.005411	1.000000	0.557509	0.759756	-0.389012	-0.053696	0.031607
subtotal	-0.002125	-0.052642	0.557509	1.000000	0.681519	0.038957	0.508594	0.130933
num_distinct_items	0.013303	-0.026407	0.759756	0.681519	1.000000	-0.445669	0.046610	0.065625
min_item_price	-0.009192	-0.041685	-0.389012	0.038957	-0.445669	1.000000	0.542757	0.042681
max_item_price	-0.005576	-0.088698	-0.053696	0.508594	0.046610	0.542757	1.000000	0.132981
total_onshift_dashers	0.069461	0.145208	0.031607	0.130933	0.065625	0.042681	0.132981	1.000000
total_busy_dashers	0.060163	0.149722	0.028508	0.125688	0.060354	0.044173	0.130909	0.943723
total_outstanding_orders	0.063385	0.135079	0.034295	0.130085	0.067517	0.041207	0.130478	0.936166
estimated_order_place_duration	-0.048696	-0.684553	-0.023991	0.034433	0.003040	0.051891	0.084114	-0.185615
estimated_store_to_consumer_driving_duration	0.017713	-0.010028	0.007443	0.038501	0.025105	0.004138	0.028656	0.045978
hour_of_day	-0.008898	0.013537	-0.070928	-0.190644	-0.119161	-0.051391	-0.188972	-0.374238
weekday	0.000287	0.000454	0.020701	0.029330	0.028227	-0.002020	0.026088	0.102843
total_seconds	-0.038318	-0.047118	0.078685	0.144317	0.106631	0.008554	0.088823	0.047002



# Code

```
data_cleaned3 = data_cleaned2.drop(['store_id', 'store_primary_category'], axis=1)
```

```
data_cleaned4 = data_cleaned3.drop(['num_distinct_items', 'total_onshift_dashers', 'total_busy_dashers', 'order_protocol'], axis=1)
```

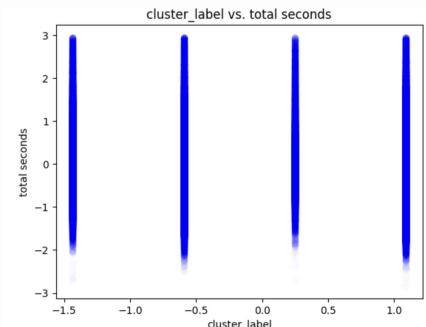
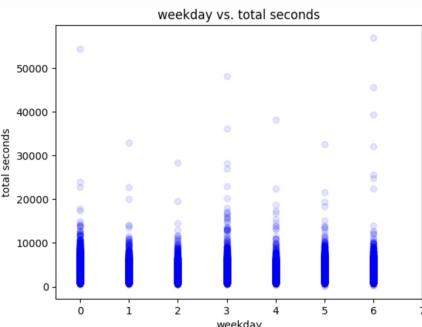
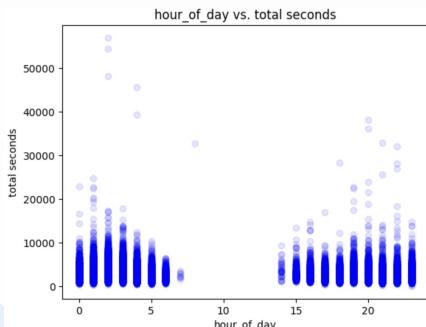
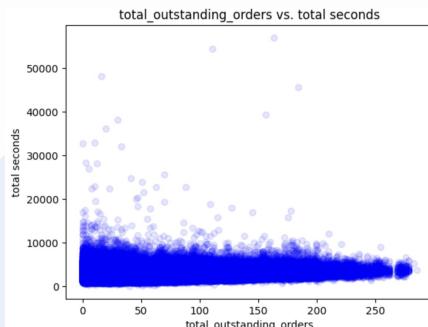
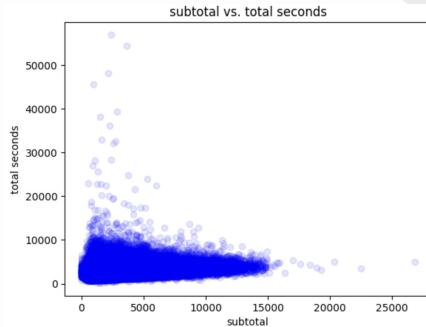
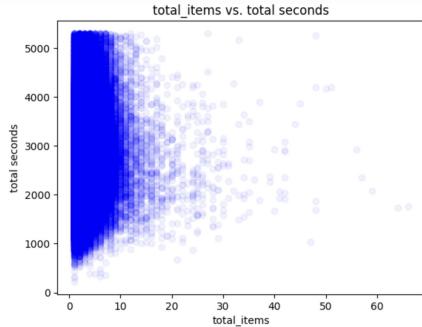
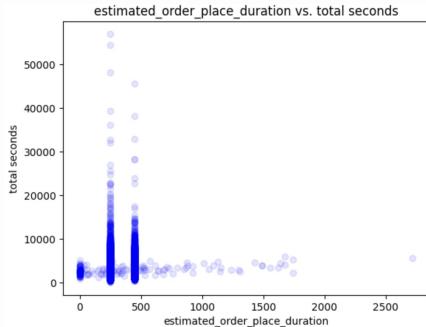
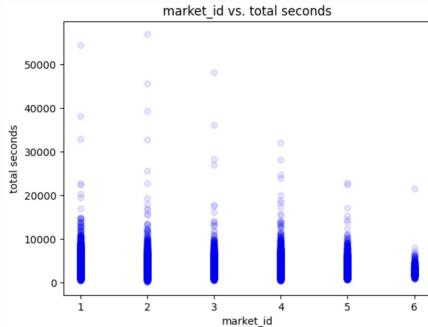


# Removing Outliers

- Plotted scatter plots comparing each feature with the response variable ('total\_seconds').
- Removed outliers for features that could have a large impact our model's results
- Removed outliers for:
  - 'total\_seconds'
  - 'subtotal'
  - 'estimated\_store\_to\_consumer\_driving\_duration'
  - 'max\_item\_price'
  - 'min\_item\_price'



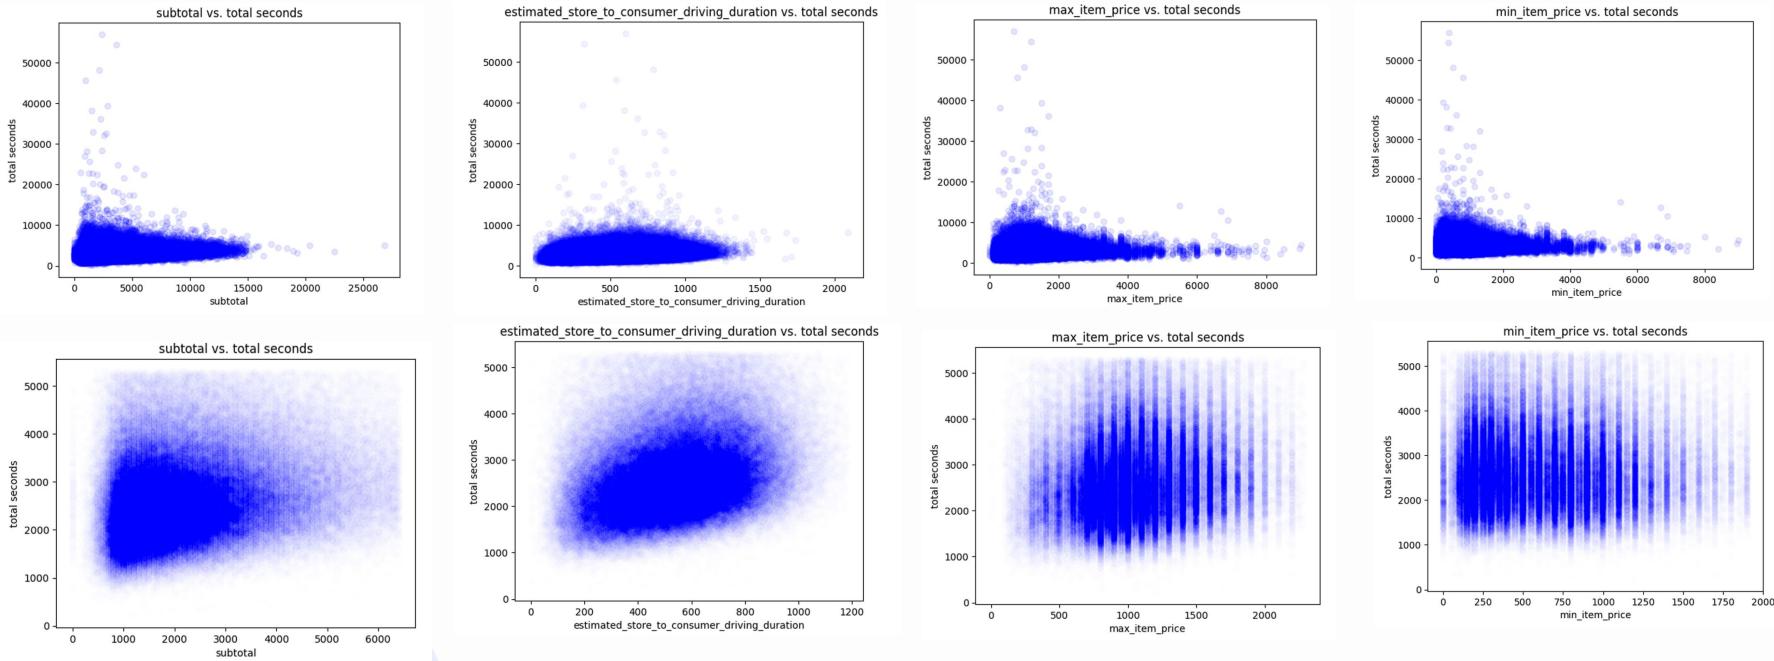
# Predictors vs. Response Scatter Plots





# Predictors vs. Response Scatter Plots

## - Before & After





# Code

```
df_temp = data_cleaned4

mask = pd.Series(True, index=df_temp.index)

columns = ['total_seconds','subtotal','estimated_store_to_consumer_driving_duration','max_item_price','min_item_price']
for column in columns:
    Q1 = df_temp[column].quantile(0.25)
    Q3 = df_temp[column].quantile(0.75)
    IQR = Q3 - Q1
    in_range = (df_temp[column] >= Q1 - 1.5 * IQR) & (df_temp[column] <= Q3 + 1.5 * IQR)
    mask &= in_range

data_no_outliers = df_temp[mask]

data_no_outliers.describe()
```



# Splitting the Data

We used the `train_test_split` function to split our data into training and testing with a test size of 20%

```
xTrain, xTest, yTrain, yTest = train_test_split(data_no_outliers.drop(['total_seconds'], axis=1),
                                                data_no_outliers['total_seconds'], test_size=0.2, random_state=42)
```



# Normalizing Data

- **Purpose:**
  - Across features, numerical range of data was widely different
    - Makes visualization of data more difficult to comprehend
    - Also impacts values of weights + biases
- **Approach:**
  - Used Z-score normalization on the features to bind them in a range

$$Z = \frac{X - \mu}{\sigma}$$



# Code

```
scaler = StandardScaler()
xTrain_scaled = pd.DataFrame(scaler.fit_transform(xTrain), columns = xTrain.columns)
xTest_scaled = pd.DataFrame(scaler.transform(xTest), columns = xTest.columns)

yTrain_scaled = pd.Series(scaler.fit_transform(yTrain.values.reshape(-1, 1)).ravel())
yTest_scaled = pd.Series(scaler.transform(yTest.values.reshape(-1, 1)).ravel())
```



02

# Unsupervised Learning



# Clustering

- K-means clustering algorithm → to understand underlying patterns in our data
- Methodology for clustering: K-means
  - Used features with clear trend when plotted against the 'total\_seconds'.
  - Used elbow method to find best K-value
    - Tried K-values 1-12
    - Best K-value = 4



# Elbow Graph Code

```
# Select K number of centroids
k_scaler = StandardScaler()
scaled_data = pd.DataFrame(k_scaler.fit_transform(data_no_outliers), columns=data_no_outliers.columns)

X_scaled = scaled_data.copy()[['total_items','subtotal','min_item_price','max_item_price']]

# Elbow method to choose K
inertia = []

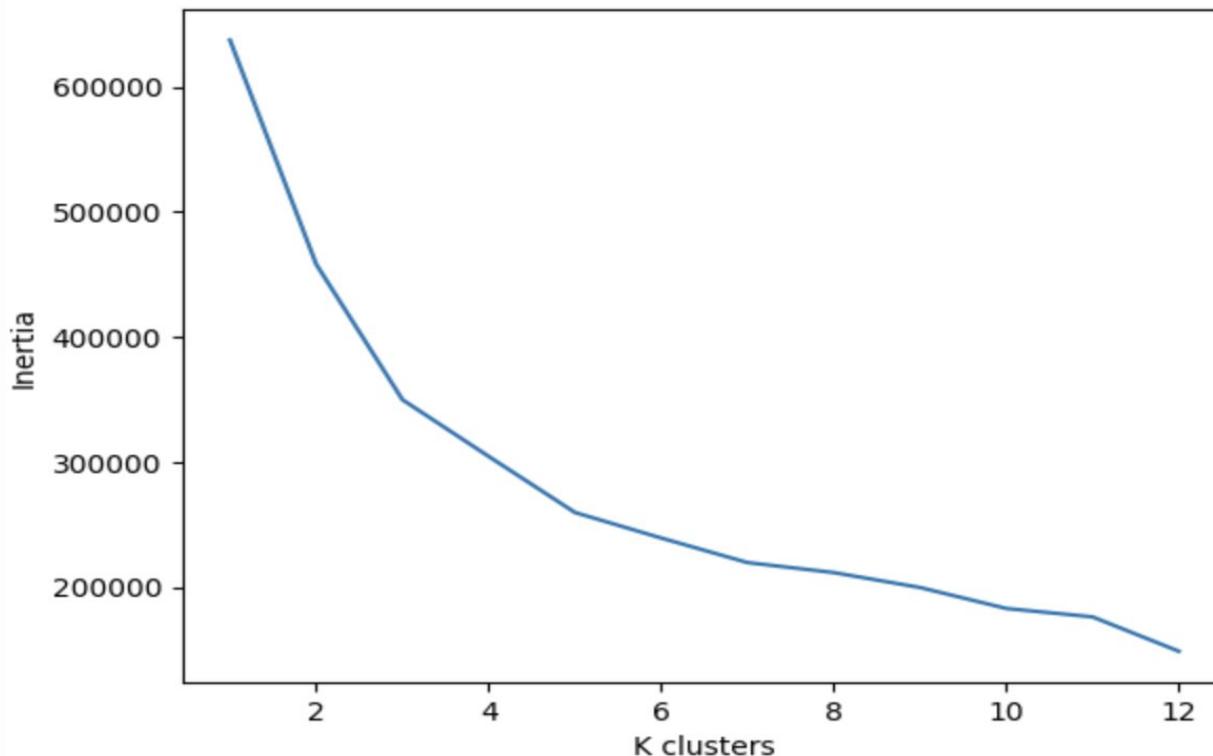
for k in range(1,13):
    kmeans = KMeans(n_clusters=k, random_state=42)
    clusters = kmeans.fit_predict(X_scaled)
    X_scaled['cluster'] = clusters
    inertia.append(round(kmeans.inertia_, 1))

plt.plot(range(1,13), inertia)
plt.xlabel('K clusters')
plt.ylabel('Inertia')
plt.title('Elbow Method for Heart Data')
plt.show()

#Selected K = 4
```

# Elbow Graph

Elbow Method for DoorDash ETA Prediction Data





# Clustering (continued)

- Application
  - We created a new feature 'cluster\_label' using the class labels obtained from clustered as values
  - Used to fit our Decision Tree and Neural Network models and make predictions

```
kmeans = KMeans(n_clusters=4)
y_kmeans = kmeans.fit_predict(X_scaled)

cluster_labels = kmeans.labels_
scaled_data['cluster_label'] = cluster_labels

print(scaled_data.groupby('cluster_label').mean())

scaled_data['cluster_label'] = k_scaler.fit_transform(scaled_data['cluster_label'].values.reshape(-1, 1))
scaled_data.head()
```



03

## Supervised Learning and Methodology

# Models constructed

**01 Linear + Polynomial Regression**

**02 Decision Trees**

**03 Neural Networks**

**04 Logistic Regression (ruled out)**

Error metrics: **5-fold Cross Validation for Mean Squared Error (MSE) and Mean Absolute Error (MAE)**





# Feature Selection - Greedy Algorithm

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

linModel = LinearRegression()
best_metric = float('inf')
selectedPredictors = []
modelTemp = []
columns = list(xTrain.columns)
```

```
for i in range(6):
    best_feature_found = None
    for column in columns:
        if column not in selectedPredictors:
            modelTemp = selectedPredictors + [column]
            linModel.fit(xTrain[modelTemp], yTrain)
            yPred = linModel.predict(xTest[modelTemp])
            mse = mean_squared_error(yTest, yPred)

            if mse < best_metric:
                best_metric = mse
                best_feature_found = column
    if best_feature_found != None:
        selectedPredictors.append(best_feature_found)
    else:
        break

print(selectedPredictors)
xTrain[selectedPredictors].head()
```

## Selected predictor variables:

'estimated\_store\_to\_consumer\_driving\_duration', 'total\_outstanding\_orders',  
'subtotal', 'estimated\_order\_place\_duration', 'hour\_of\_day', 'market\_id'



# Linear/Polynomial Regression

- Fitted data to a linear model (degree 1)
- Used polynomial regression to transform it... (degree 2, 3, 4 and 5)
- Applied regularization; MSE and MAE similar to original.
- Used cross-validation; MSE and MAE slightly better.
- Degree 4 is the best polynomial with the lowest MSE and MAE
  - **MSE = 0.8065**
  - **MAE = 0.7120**

Degree 1

Original mse: 0.8325774853912852  
Original mae: 0.7258654654301407  
Ridge mse: 0.8325759910660239  
Ridge mae: 0.7258707557910781  
cv mse: 0.8387657458135488  
cv mae: 0.7282163885849909

Degree 2

Original mse: 0.8132135994526223  
Original mae: 0.7167430035193708  
Ridge mse: 0.8132099844833449  
Ridge mae: 0.716748915130252  
cv mse: 0.8183351693750502  
cv mae: 0.718362195314951

Degree 3

Original mse: 0.8046651597150392  
Original mae: 0.7129076916913976  
Ridge mse: 0.804666798175454  
Ridge mae: 0.7129113725097594  
cv mse: 0.8100373143909086  
cv mae: 0.7149839917191115

Degree 4

Original mse: 0.7991966753348451  
Original mae: 0.7102773757185288  
Ridge mse: 0.7990076408304524  
Ridge mae: 0.7102587901815901  
cv mse: 0.8065482230008808  
cv mae: 0.7120245876834508

Degree 5

Original mse: 6.629469077441348  
Original mae: 0.7166952463406273  
Ridge mse: 0.8270773057130727  
Ridge mae: 0.7040410389937369  
cv mse: 11.063890741220808  
cv mae: 0.7222488319893449

# Linear/Polynomial Regression (code)

```
for degree in degrees:
    print(f"Degree {degree}")
    # Original Polynomial Regression
    polyreg = make_pipeline(PolynomialFeatures(degree=degree), LinearRegression())
    polyreg.fit(xTrain_scaled[selectedPredictors], yTrain_scaled)
    y_pred = polyreg.predict(xTest_scaled[selectedPredictors])
    # Original MSE and MAE
    mse = mean_squared_error(yTest_scaled, y_pred)
    mae = mean_absolute_error(yTest_scaled, y_pred)
    print(f"Original mse: {mse}")
    print(f"Original mae: {mae}")

# Grid search for best alpha
best_alpha = 1.0
best_cv_mse = np.inf
for alpha in alphas:
    polyreg_ridge = make_pipeline(PolynomialFeatures(degree=alpha), Ridge(alpha=alpha))
    cv_mse = -cross_val_score(polyreg_ridge, xTrain_scaled[selectedPredictors], yTrain_scaled, cv=5)
    if cv_mse < best_cv_mse:
        best_cv_mse = cv_mse
        best_alpha = alpha
# Polynomial Regression with Ridge Regularization
polyreg_ridge = make_pipeline(PolynomialFeatures(degree=degree), Ridge(alpha=best_alpha))
polyreg_ridge.fit(xTrain_scaled[selectedPredictors], yTrain_scaled)
y_pred_ridge = polyreg_ridge.predict(xTest_scaled[selectedPredictors])
# Ridge MSE and MAE
mse_ridge = mean_squared_error(yTest_scaled, y_pred_ridge)
mae_ridge = mean_absolute_error(yTest_scaled, y_pred_ridge)
print(f"Ridge mse: {mse_ridge}")
print(f"Ridge mae: {mae_ridge}")
# Cross-Validated MSE and MAE
cv_scores_mse = cross_val_score(polyreg_ridge, xTrain_scaled[selectedPredictors], yTrain_scaled, cv=5)
cv_mse = -cv_scores_mse.mean()
cv_scores_mae = cross_val_score(polyreg_ridge, xTrain_scaled[selectedPredictors], yTrain_scaled, cv=5)
cv_mae = -cv_scores_mae.mean()
print(f"cv mse: {cv_mse}")
print(f"cv mae: {cv_mae}")
```



# Decision Tree

- Deciding on hyperparameters before fitting the data...
  - Implemented a nested for loop to minimize MSE for different hyperparameters
    - Max depth values tested: 10–50
    - Min samples values tested: 100–500 with step = 100
  - Final max depth value: 15
  - Final minimum samples per leaf: 200
- **MSE = 0.7975**
- **MAE = 0.707**





# Random Forest

- Hyperparameters
  - Originally started with the same hyperparameters as the decision tree
  - Looked through documentation and played around with other attributes of the object (impurity decrease, bootstrap, warm start, etc.)
    - Didn't perform better than the singular decision tree :(
  - Dropped limitations, set min samples leaf as 50 and left everything else as default
- **MSE = 0.766**
- **MAE = 0.6929**



# Decision Tree (code)

## Finding best hyperparameters

```
1 decTreeMSEs = []
2 decTreeValuesTested = []
3
4 for i in range(10, 51):
5     for j in range(100, 501, 100):
6         decTreeIter = DecisionTreeRegressor(max_depth = i, min_samples_leaf = j, random_state = 0)
7         decTreeIter.fit(xTrain_scaled, yTrain_scaled)
8         decTreeIterPreds = decTreeIter.predict(xTest_scaled)
9         decTreeIterMSE = mean_squared_error(yTest_scaled, decTreeIterPreds)
10        decTreeValuesTested += [(i, j)]
11        decTreeMSEs += [decTreeIterMSE]
12
13 decTreeMinMSE = min(decTreeMSEs)
14 decTreeMinMSEIndex = decTreeMSEs.index(decTreeMinMSE)
15 print(decTreeMinMSE, decTreeValuesTested[decTreeMinMSEIndex])
0.8007229225962674 (15, 200)
```

## Decision Tree Fitting & MSE

```
1 # Decision Tree
2 from sklearn.tree import DecisionTreeRegressor
3
4 decisionTree = DecisionTreeRegressor(max_depth = 15, min_samples_leaf = 200, random_state=0)
5
6 decisionTree.fit(xTrain_scaled, yTrain_scaled)
7 preds = decisionTree.predict(xTest_scaled)
8
9 kfold = model_selection.KFold(n_splits=10, random_state=42, shuffle=True)
10 cvScoreDT = model_selection.cross_val_score(decisionTree, xTrain_scaled, yTrain_scaled, scoring = 'neg_mean_squared_error', cv=kfold)
11 cvMaeDT = model_selection.cross_val_score(decisionTree, xTrain_scaled, yTrain_scaled, scoring = 'neg_mean_absolute_error', cv=kfold)
```

```
1 print(f"MSE for Decision Tree: {(-1 * cvScoreDT).mean()}")
2 print(f"MAE for Decision Tree: {(-1 * cvMaeDT).mean()}")
```

```
MSE for Decision Tree: 0.7959110316323239
MAE for Decision Tree: 0.7070023374030607
```

# Random Forest (code)

```
1 # Random Forest
2
3 randomForest = RandomForestRegressor(n_estimators = 100, min_samples_leaf = 50, random_state = 0)
4 randomForest.fit(xTrain_scaled, yTrain_scaled)
5 randomForestPred = randomForest.predict(xTest_scaled)
6 randomForestScore = mean_squared_error(yTest_scaled, randomForestPred)
7
8 kfold = model_selection.KFold(n_splits=10, random_state=42, shuffle=True)
9 cvScoreRF = model_selection.cross_val_score(randomForest, xTrain_scaled, yTrain_scaled, scoring = 'neg_mean_squared_error', cv=kfold)
10 cvMaeRF = model_selection.cross_val_score(randomForest, xTrain_scaled, yTrain_scaled, scoring = 'neg_mean_absolute_error', cv=kfold)

1 print(f"MSE for Random Forest: {(-1 * cvScoreRF).mean()}")
2 print(f"MAE for Random Forest: {(-1 * cvMaeRF).mean()}")
```

MSE for Random Forest: 0.7645830826463217

MAE for Random Forest: 0.6929210387572369



# Neural Networks

- Wanted a more sophisticated model that would (hopefully) capture relationships between predictor and response variables better.
- All predictor variables were used to create the most accurate model possible
- Hyperparameters used:
  - 3 hidden layers with 64, 32, & 16 nodes respectively
  - Activation function = ReLU
  - Lambda for L2 regularization = 0.01
  - 200 epochs
- Results:
  - **MSE = 0.7605**
  - **MAE = 0.6877**



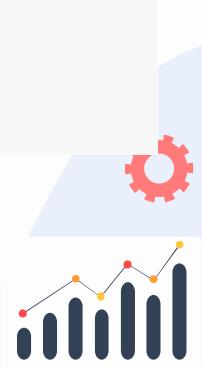


# Neural Network - Regularization Term

```
from sklearn.neural_network import MLPRegressor
lambdaList = [0.001, 0.01, 0.1, 1, 5, 10, 20, 100, 1000]

for l in lambdaList:
    NNmodel = MLPRegressor(hidden_layer_sizes=(64,32,16), activation='relu',
                           alpha=l, solver="adam", max_iter=200, random_state=42)

    NNmodel.fit(xTrain_scaled, yTrain_scaled)
    y_pred = NNmodel.predict(xTest_scaled)
    mse = mean_squared_error(yTest_scaled, y_pred)
    print(mse)
```





# Neural Network - Regularization Term

0.01 lambda value

0.7555421875751395  
0.745209679982672  
0.749605055882063  
0.8013635315322923  
0.8316754300481609  
0.8416023234924026  
0.8970802038477983  
0.9896966936187183  
0.9896978446110412





# Neural Networks (code)

```
kfold = model_selection.KFold(n_splits=10, random_state=42, shuffle=True)
NNmodel = MLPRegressor(hidden_layer_sizes=(64,32,16), activation='relu',
                       alpha=0.1, solver="adam", max_iter=500, random_state=42)

cvScoreNNmse = model_selection.cross_val_score(NNmodel, xTrain_scaled, yTrain_scaled, scoring = 'neg_mean_squared_error', cv=kfold)
cvScoreNNmae = model_selection.cross_val_score(NNmodel, xTrain_scaled, yTrain_scaled, scoring = 'neg_mean_absolute_error', cv=kfold)

print(f"CV Score (MSE): {-1 * cvScoreNNmse}")
print(f"CV Score (MAE): {-1 * cvScoreNNmae}")
print("")

CV Score (MSE): [0.76493588 0.76572064 0.74217203 0.7560194 0.75387734 0.77964625
 0.77026689 0.77669196 0.8198754 0.74715611]
CV Score (MAE): [0.69121065 0.68810007 0.67996133 0.69268253 0.6962632 0.72870861
 0.6846468 0.67966756 0.69742232 0.68261918]
```

```
print(f"CV-score (MSE) for Neural Network: {(-1 * cvScoreNNmse).mean()}")
print(f"CV-score (MAE) for Neural Network: {(-1 * cvScoreNNmae).mean()}"
```

```
CV-score (MSE) for Neural Network: 0.7605185310574005
CV-score (MAE) for Neural Network: 0.6877246747009738
```





# Logistic Regression

- Did **NOT** end up considering logistic regression for final model
- Code on how a logistic regression model could've worked...

```
# Logistic Regression

from sklearn.model_selection import KFold
from sklearn.linear_model import LogisticRegression

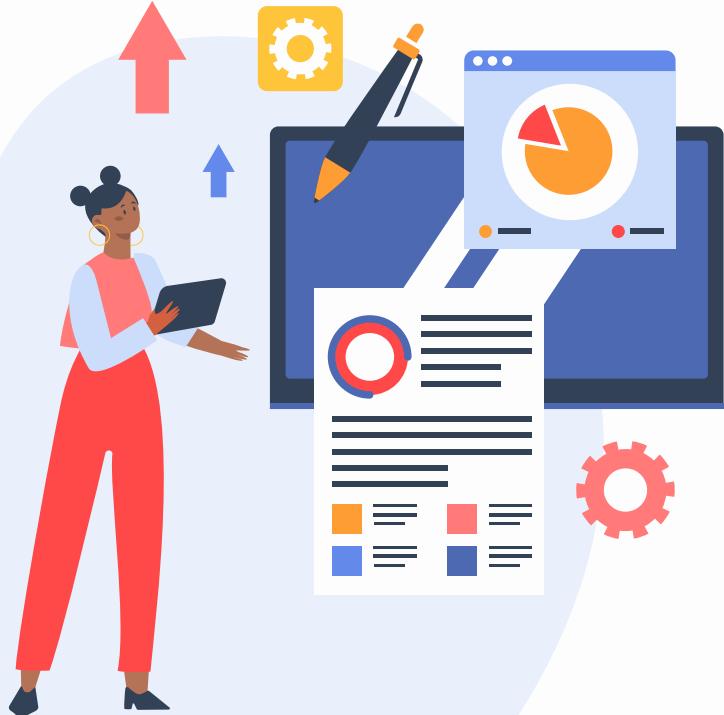
logRegModel = LogisticRegression()
logRegModel.fit(xTrain[selectedPredictors], yTrain)
logRegPred = logRegModel.predict(xTest)
logRegKFold = KFold(n_splits = 10, random_state = 5, shuffle = True)
logRegKFoldScore = model_selection.cross_val_score(logRegModel, xTrain[selectedPredictors], yTrain, cv = 5)
print(f"Logistic Regression mean kfold score: {logRegKFoldScore.mean()}")
```





04

## Presentation of Results



WHICH  
MODEL DID  
WE  
CHOOSE?

# NEURAL NETWORKS

...because it had the best MSE and MAE of all our models



# Full List of Error Metrics

## Linear Regression

MSE = 0.8390

MAE = 0.728

## Polynomial regression

MSE = 0.8073

MAE = 0.7120

## Decision Tree

MSE = 0.7974

MAE = 0.7070

## Random Forest

MSE = 0.7660

MAE = 0.6929

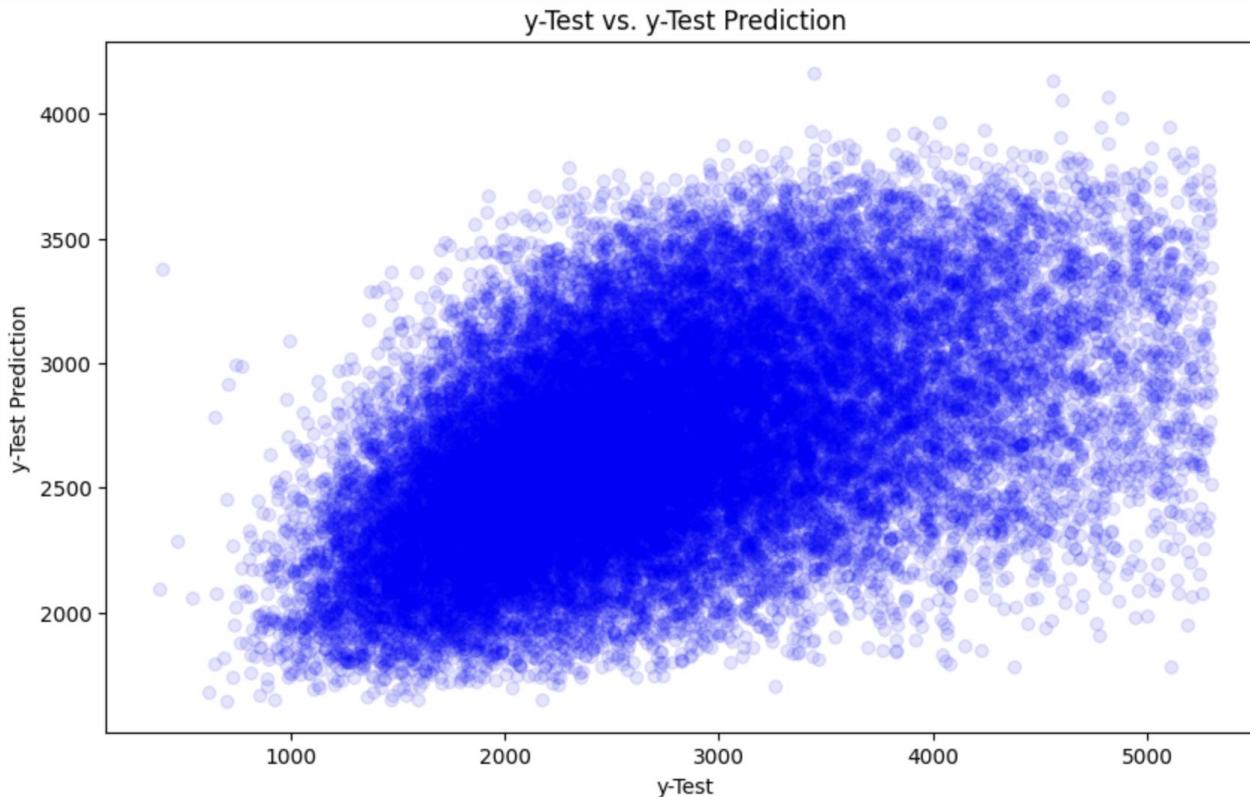
## Neural Network

MSE = 0.7605

MAE = 0.6877



# Y-Test vs. Y-Test Prediction Graph





# Summary

- Neural Network Model was used to predict the **total delivery time** of a DoorDash order in seconds
- Used all the predictors (pre-existing and those we created) to generate the best model.
- Why did we choose Neural Networks?
  - Universal function approximators
  - Can discover non-linear relationships to create more accurate predictions.



# **THANK YOU!!!**

**(questions?)**

